The Architecture of Spuriousness: Investigating Shortcut Learning and Generalization Failures in Large-Scale Neural Models

The persistent success of neural networks on standardized benchmarks has long been contrasted with their fragile generalization in the face of out-of-distribution (OOD) shifts. This divergence is fundamentally rooted in the phenomenon of shortcut learning—a behavioral failure where models identify and exploit spurious correlations within training datasets rather than internalizing the intended causal or semantic structures.[1] Historically, these failures were characterized in linear models and shallow settings as the "Clever Hans" effect, named after the horse that appeared to solve mathematical problems but was actually responding to the subtle cues of his observers.[3] As machine learning has evolved from simple task-specific algorithms to massive, transformer-based Large Language Models (LLMs), the mechanisms through which these biases propagate and manifest have become increasingly complex.[3] Understanding the trajectory from linear inductive biases to the multi-layered representational hierarchies of LLMs is essential for developing robust artificial intelligence capable of reliable performance in real-world scenarios.[1]

Theoretical Foundations and the Linear Lineage of Shortcut Bias

The conceptual core of shortcut learning lies in the mismatch between intended and learned solutions. In the standard Empirical Risk Minimization (ERM) paradigm, a model is trained to minimize a loss function on a finite dataset, which often contains artifacts—statistical associations that happen to correlate with the labels but do not represent the underlying task logic.[6] These non-robust features offer a "path of least resistance" for the optimization process, allowing the model to achieve high accuracy on Independent and Identically Distributed (IID) test sets while remaining oblivious to the actual semantic content.[6]

Inductive Bias and Function Complexity

Generalization in neural networks is largely governed by their inductive biases, which determine the types of functions the models are likely to represent. Research into the "Neural Redshift" suggests that neural architectures are not neutral but possess a strong bias toward functions of lower complexity, a property often referred to as simplicity bias.[5] This bias is not inherent to the concept of deep learning but is a direct consequence of specific architectural building blocks. Components such as Rectified Linear Units (ReLUs), residual connections, and layer normalization induce a preference for functions with low frequency, low order, and high compressibility.[5]

In linear models, this bias is straightforward: the model can only learn linear relationships. However, deep networks, despite their vast representational capacity, often behave like linear systems during the early stages of training.[9] The inclusion of bias terms in ReLU networks plays a pivotal role in this process, pushing the units into a transient linear regime where the model learns an "optimal constant solution" (OCS) that mirrors the distribution of target labels while ignoring input information.[10] This early linear phase reinforces a low-rank bias, encouraging the model to adopt simple decision rules—the very shortcuts that lead to generalization failures when the OOD data deviates from the training set's spurious patterns.[9]

Taxonomy of Spurious Correlations

The variety of dataset biases can be categorized according to their manifestations in natural language and vision tasks. These shortcuts range from low-level lexical cues to higher-level conceptual associations.

Shortcut Type	Definition and Mechanism	Example in Natural Language Processing
Lexical Bias	High statistical correlation between specific tokens (e.g., stop words, negation) and class labels.	Predicting "contradiction" every time the word "never" or "no" appears in a hypothesis.[6]
Overlap Bias	Reliance on the degree of shared features or n-grams between two input branches (e.g., premise and hypothesis).	Assuming "entailment" solely because the hypothesis is composed of words from the premise.[11]
Position Bias	Use of positional cues or the skewness of answer locations within a training set.	Prioritizing the beginning of a sentence while neglecting the semantic conclusion at the end.[1]
Style Bias	Distinguishing between classes based on stylistic markers (e.g., tone,	A sentiment analyzer associating formal language with negative

	formality, length) rather than content.	reviews due to training set imbalances.[6]
Concept Bias	Dependence on higher-level categories or entities that are spuriously correlated with specific outcomes.	Associating specific demographic entities or "popular" mentions with a specific label regardless of the context.[13]

Propagation Mechanisms in Transformer Architectures

While transformers have revolutionized sequential modeling, they are not immune to the simplicity biases identified in earlier models. In fact, they inherit and amplify these biases through their unique structural components. The attention mechanism, while allowing for complex global dependencies, also provides a mechanism for models to "attend" to spurious markers with high precision.[3]

Architectural Robustness vs. Data Sensitivity

Research comparing transformers to older n-gram models reveals a nuanced picture of bias propagation. While n-gram models are highly sensitive to context window size—where larger windows often exacerbate the capturing of spurious local patterns—transformers demonstrate a form of architectural robustness.[15] The bias scores in transformers remain relatively stable regardless of the number of layers or attention heads, suggesting that the propagation of bias is less a function of depth and more a consequence of the training data's quality and the initial pre-training objectives.[15]

However, the "Neural Redshift" is highly relevant to transformer sequence models. Random-weight transformers, even before training, produce sequences of low complexity, inheriting this prior from the same building blocks that drive simplicity in Multi-Layer Perceptrons (MLPs).[5] The interaction between components like Gaussian Error Linear Units (GELUs) and residual connections creates a representational hierarchy that favors simple pattern matching over complex logical reasoning, especially during the early phases of pre-training.[8]

Sensitivity and Robustness Metrics

The inductive bias of transformers can be quantified through their sensitivity to token-wise random perturbations. Transformers generally exhibit lower sensitivity than MLPs, CNNs, and LSTMs across both vision and language tasks.[17] This low-sensitivity bias has significant implications:

1. **Improved Robustness**: Lower sensitivity correlates with better handling of noise, yet it also reflects a flatter loss landscape that might settle into simpler, shortcut-based minima.[17]
2. **Positional Anchoring**: Transformers are particularly sensitive to specific "anchor" tokens, such as the `` token, which often acts as a aggregator for both robust semantic features and spurious dataset artifacts.[17]
3. **Token Distribution Mimicry**: LLMs often exhibit "lazy" behavior by mimicking prompt formats and token distributions, leading to hallucinations when the expected surface-level patterns shift.[3]

Systematic Evaluation: The Shortcut Suite and LLM Vulnerabilities

To rigorously evaluate the susceptibility of state-of-the-art LLMs to shortcut learning, researchers have developed comprehensive evaluation frameworks like the Shortcut Suite. This test suite focuses on six primary shortcut types—Lexical Overlap, Subsequence, Constituent, Negation, Position, and Style—across various Natural Language Inference (NLI) and understanding tasks.[1]

Key Findings in LLM Shortcut Performance

Extensive experiments on LLMs, including both closed-source and open-source models, have yielded several critical insights into their robustness.

Evaluation Metric	Observed Behavior in Shortcut-Laden Datasets	Performance Impact
Accuracy	Significant drops across all shortcut types, particularly Constituent and Negation.	Drops of over 40% in some cases when shortcuts are removed or neutralized.[1]
Confidence	LLMs consistently exhibit overconfidence in their predictions when utilizing shortcuts.	Models report high probability for incorrect labels if a shortcut is present.[1]
Explanation Quality	Lower quality in shortcut-laden datasets compared to standard	Errors categorized into distraction, disguised comprehension, and logical fallacies.[1]

	benchmarks.	
Prompt Sensitivity	Zero-shot and few-shot prompts often exacerbate the reliance on surface cues.	Few-shot scenarios sometimes perform worse than zero-shot due to captured spurious correlations in examples.[1]

The evaluation highlights that even the most capable models falter on the Position dataset, demonstrating a propensity to prioritize the beginning of sentences while neglecting crucial information at the end.[1] This vulnerability suggests that the sequential nature of autoregressive generation reinforces a "first-impression" bias that ignores the full semantic context.

The Role of Prompting Strategies

The method of interacting with LLMs significantly influences their reliance on shortcuts. Chain-of-Thought (CoT) prompting has emerged as a particularly effective, albeit imperfect, mitigation strategy.[1] By requiring the model to generate intermediate reasoning steps, CoT forces a level of semantic processing that simple zero-shot prompting avoids. However, even within CoT traces, models exhibit "disguised comprehension," where the reasoning appears logically sound but ultimately leads to a pre-determined, shortcut-based answer.[1]

The Scaling Paradox: Why Larger Models are "Lazier Learners"

One of the most counterintuitive findings in the study of large-scale models is the relationship between model capacity and shortcut susceptibility. While one might expect that more parameters would lead to better feature extraction and a deeper understanding of semantics, the opposite is often observed: larger models are frequently more prone to shortcut learning than smaller ones.[1]

Reverse Scaling of Spurious Correlations

This "reverse scaling" phenomenon is attributed to the increased capacity of larger models to identify and exploit complex statistical patterns. While a small model might lack the representational power to find a subtle shortcut, a larger model is highly efficient at discovering any mathematical correlation that reduces the training loss.[1] LLMs are described as "lazy learners" because they prioritize these easy-to-learn mappings over the high-level semantic reasoning required for true generalization.[21]

Experimental data on tasks like SST2 and various information extraction benchmarks shows that as model size increases (e.g., from OPT-2.7B to OPT-13B), the performance drop on anti-shortcut test sets becomes more pronounced.[21] Larger models receive a more significant penalty because they have committed more heavily to the spurious triggers—such as signs, common words, or specific text styles—that were present in the in-context learning prompts.[21]

Strategic Selection vs. Knowledge Acquisition

A deeper analysis of model sizes (7B–8B vs. 32B+ parameters) suggests that the superior performance of larger models on standard benchmarks is not necessarily due to having "more accurate" knowledge but rather a superior "strategy selection".[13] Larger models can selectively rely on robust numerical knowledge when it is perceived as highly reliable, but they fall back on heuristics (like entity popularity or mention order) more effectively than smaller models.[13] In contrast, smaller models often show no such discrimination, relying on shortcuts even when they possess the underlying factual knowledge to answer correctly.[13]

Optimization and Training Dynamics: The Small Batch Advantage

The choice of optimization strategy, particularly the batch size and the specific optimizer configuration, has a profound impact on a model's tendency to settle into shortcut-based solutions. Conventional wisdom has long dictated that large batch sizes are essential for the stability and efficiency of language model pre-training. However, recent research challenges this assumption, suggesting that gradient accumulation—the standard method for achieving large effective batch sizes—is often "wasteful" and potentially detrimental to robustness.[23]

Robustness to Hyperparameters

Small batch sizes, reaching all the way down to batch size one, have been found to offer several training advantages that directly combat the instability and sensitivity typically associated with LLM training.[23]

- **Stable Training**: Small batch sizes train stably and are consistently more robust to hyperparameter choices.[23]
- **Broad Low-Loss Regions**: For batch sizes as small as one, the model achieves nearly optimal loss across a wide range of learning rates and momentum values.[24] This is in sharp contrast to large batch sizes (e.g., 512+), which are highly sensitive to these values because they require massive learning rates to compensate for fewer optimization steps.[23]
- **Competitive Per-FLOP Performance**: Smaller batch sizes achieve equal or better per-FLOP performance because they allow for more frequent updates. Training often occurs in a "far-from-convergence" regime where frequent, noisier updates are more efficient than large, infrequent steps.[23]

Vanilla SGD and Optimizer Scaling Rules

Remarkably, small batch regimes enable stable language model training using vanilla Stochastic Gradient Descent (SGD) without momentum or weight decay.[23] This finding is significant because it eliminates the need for storing optimizer states, providing a memory footprint similar to Low-Rank Adaptation (LoRA) while maintaining the benefits of full fine-tuning.[23]

When using adaptive optimizers like Adam at small batch sizes, a new scaling rule is proposed to maintain stability. Rather than holding the decay rate of the second moment (β_2) fixed, researchers recommend holding its **half-life** fixed in terms of tokens.[25] This adjustment ensures that the optimizer's memory of past gradient noise remains consistent relative to the data seen, preventing the "vanishing" of gradients as the batch size scales down.[23]

Optimization Parameter	Traditional Large Batch Approach	Proposed Small Batch Heuristic
Learning Rate	High, requires careful tuning.	Lower, more robust across a wide range.
Optimizer	Sophisticated (Adam, Muon).	Simple (Vanilla SGD).
β_2 (Second Moment)	Fixed decay rate across steps.	Fixed half-life in terms of tokens.
Memory Overhead	High due to optimizer states.	Minimal (zero optimizer state for vanilla SGD).
Throughput	Maximized via gradient accumulation.	Maximized by smallest batch that saturates MFU.

The Epochal Sawtooth Phenomenon (ESP)

A recurring training loss pattern known as the Epochal Sawtooth Phenomenon (ESP) has been identified in adaptive gradient-based optimizers like Adam.[27] ESP is characterized by a spike in loss at the start of each epoch, followed by a decline. This pattern is driven by:

1. **Adaptive Learning Rate Adjustments**: The second moment estimate reacts to the reshuffling of data at epoch boundaries.[27]
2. **Immediate Re-exposure Effect**: During data shuffling, the model may encounter the same samples in rapid succession, leading to a form of localized "memorization" or shortcut learning at the beginning of each epoch.[27]

3. **Regularization via β_2**: Smaller values of β_2 tend to exacerbate ESP, but they can also act as a form of regularization, preventing the model from becoming too rigid in its learned patterns.[27]

Advanced Mitigation Strategies and Robust Optimization Frameworks

Given the prevalence of shortcut learning, several advanced frameworks have been developed to detect, diagnose, and mitigate these behaviors. These methods range from data-centric augmentation to model-centric regularization and structural guidance.

COMI and MiMu: Multi-Shortcut Mitigation

The **COrrect and MItigate (COMI)** strategy addresses the over-reliance on shortcuts by integrating information separation with standard Empirical Risk Minimization.[1] COMI encourages models to focus on "real" features by penalizing the capture of spurious correlations through a random masking strategy.[29]

Building on this, the **MiMu** method is designed to mitigate *multiple* unknown shortcuts in real-world scenarios where cues are diverse and unknown.[30]

- **Self-Calibration Strategy**: Implemented in a source model to prevent the model from relying on shortcuts and making overconfident predictions. It aligns the model's prediction with its confidence through a specific calibration loss.[30]
- **Self-Improvement Strategy**: Implemented in a target model to further reduce reliance on multiple shortcuts. It uses a **Random Mask Strategy** to obscure partial attention positions, forcing the model to diversify its focus rather than fixating on a single, fixed shortcut region.[30]
- **Adaptive Attention Alignment**: The target model's attention weights are aligned with those of the calibrated source model, guiding the target model toward more robust features without requiring additional manual supervision or post-hoc attention map generation.[29]

Smart Data Augmentation: SAug-LLM

Traditional data augmentation often requires manual definitions of shortcuts, which limits its scalability. **SAug-LLM** (Smart Data Augmentation based on LLMs) leverages the autonomous capabilities of frontier models to identify and neutralize shortcuts.[33]

1. **Autonomous Discovery**: An LLM analyzes training data to identify both explicit and implicit shortcuts.[33]
2. **Anti-Shortcut Generation**: The LLM rephrases the samples to render the identified shortcuts ineffective while ensuring the labels remain valid.[33]

3. **Dual Validation**: To mitigate "confirmation bias"—where a model tends to ignore new data that contradicts its prior, biased beliefs—SAug-LLM employs a dual validation strategy during retraining.[33]
4. **Efficiency Optimization**: The augmentation process is optimized to effectively rectify biases while minimizing the consumption of additional data, achieving a performance improvement of approximately 5.61% across various NLP tasks.[33]

SSR: Structural Skeleton-guided Reasoning

In reasoning-intensive tasks, LLMs often suffer from "post-hoc rationalization," where they anchor their entire explanation to a pre-determined (often shortcut-based) response.[37] **Structural Skeleton-guided Reasoning (SSR)** breaks this cycle by decoupling the structure of reasoning from its content.[37]

- **Answer-Invariant Skeleton**: The model first generates a skeleton of functional tags (e.g., PLAN $\rightarrow$ INFER $\rightarrow$ CONCLUDE) that describes the reasoning path without referencing the specific answer.[1]
- **Content-Neutral Guidance**: This skeleton is then used to guide the generation of the full reasoning chain. By redirecting information flow toward structural planning, SSR reduces "probabilistic anchoring" to shortcuts.[37]
- **Distilled SSR (SSR-D)**: Models can be fine-tuned on teacher-generated SSR traces to produce unanchored chains that achieve up to a 10% improvement over standard suppression baselines while preserving OOD generalization.[37]

Diagnostics and Robustness Metrics for Future AI Development

Traditional evaluation metrics like IID accuracy fail to capture the underlying fragility of models relying on shortcuts. A more comprehensive approach requires metrics that assess explainability, consistency, and causal grounding.

Novel Metrics for LLM Evaluation

The Shortcut Suite and related research suggest several new dimensions for evaluating robustness.

Metric	Definition	Purpose in Robustness Evaluation
Semantic Fidelity Score	Measures the alignment between the model's	Detects "right for the wrong reasons" scenarios where

(SFS)	reasoning trace and the ground-truth semantic requirements.	the answer is correct but the logic is flawed.[1]
Internal Consistency Score (ICS)	Evaluates whether the model's explanations are consistent across different variations of the same prompt.	Identifies if a model is flip-flopping based on irrelevant surface cues.[1]
Deceptive Signal Metric	Quantifies the shift in model behavior when sensitive shortcut information is withheld during training.	Isolates the degree of reliance on task-irrelevant patterns.[29]
Spurious Correctness (S-Corr)	The fraction of label-correct decisions with rationales misaligned with golden judgments.	Detects "deceptive alignment" in generative reward models.[38]

Causal Inference and Counterfactual Auditing

As LLMs are integrated into high-stakes domains like healthcare and law, the need for causal grounding becomes paramount. Shortcut learning in medical AI, for instance, can overestimate model performance by up to 20% due to "data acquisition biases" (DAB)—spurious correlations between scanner types or hospital-specific protocols and patient outcomes.[39]

To mitigate this, researchers propose:

- **Counterfactual Node Generation**: Using adversarial networks to create "what-if" scenarios that challenge the model to generalize across different sensitive attributes.[40]
- **Multi-relational Causal Intervention (MMCI)**: Disentangling causal features from shortcut features by modeling inputs as multi-modal causal graphs.[41]
- **Subgroup Analysis**: Comparing model performance on specific data slices to identify where spurious correlations dominate.[3]

The Horizon of Large-Scale Model Generalization

The investigation into shortcut learning in large-scale models reveals a transition from simple lexical artifacts to complex structural and strategic biases. While scaling has provided unprecedented capabilities, it has also amplified the "lazy learning" tendencies of neural networks, where the pursuit of training efficiency often comes at the cost of semantic depth and logical rigor.[21]

The move toward robust artificial intelligence necessitates a departure from standard Empirical Risk Minimization. Strategies such as small batch size optimization, fixed token-based hyperparameter scaling, and autonomous anti-shortcut data generation provide a foundation for more resilient models.[23] Furthermore, by incorporating structural skeletons and causal interventions, researchers can steer LLMs away from the "Clever Hans" traps that currently undermine their reliability in the wild.[3] Ultimately, the goal of future research is to align the inductive biases of our architectures with the fundamental causal structures of human language and the physical world, ensuring that "intelligence" is not merely an illusion of statistical pattern matching.[3]

Works cited

1. Do LLMs Overcome Shortcut Learning? An Evaluation of ... - arXiv.org, accessed on February 23, 2026, https://arxiv.org/html/2410.13343v1
2. Do LLMs Overcome Shortcut Learning? An ... - ACL Anthology, accessed on February 23, 2026, https://aclanthology.org/2024.emnlp-main.679.pdf
3. Unmasking the Clever Hans effect in AI models: shortcut learning, accessed on February 23, 2026, https://www.frontiersin.org/journals/artificial-intelligence/articles/10.3389/frai.2025.1692454/full
4. Shortcut Learning in Deep Learning | by samuel chazy, accessed on February 23, 2026, https://ai.plainenglish.io/shortcut-learning-in-deep-learning-1db628968c91
5. Neural Redshift: Random Networks are not Random Functions, accessed on February 23, 2026, https://openaccess.thecvf.com/content/CVPR2024/papers/Teney_Neural_Redshift_Random_Networks_are_not_Random_Functions_CVPR_2024_paper.pdf
6. Shortcut Learning of Large Language Models in Natural ... - arXiv, accessed on February 23, 2026, https://arxiv.org/pdf/2208.11857
7. Shortcut Learning of Large Language Models in Natural Language, accessed on February 23, 2026, https://cacm.acm.org/research/shortcut-learning-of-large-language-models-in-natural-language-understanding/
8. Neural Redshift: Random Networks are not Random Functions, accessed on February 23, 2026, https://cvpr.thecvf.com/virtual/2024/poster/29922
9. Bias-mediated early linearization drives simplicity bias in ReLU, accessed on February 23, 2026, https://openreview.net/forum?id=oqP3qkHAsA
10. Simplicity Bias in Transformers and their Ability to Learn Sparse, accessed on February 23, 2026, https://www.researchgate.net/publication/372918730_Simplicity_Bias_in_Transformers_and_their_Ability_to_Learn_Sparse_Boolean_Functions
11. yyhappier/ShortcutSuite - GitHub, accessed on February 23, 2026, https://github.com/yyhappier/ShortcutSuite

12. A Comprehensive Analysis of Shortcut Learning in Text ... - arXiv, accessed on February 23, 2026, https://arxiv.org/html/2409.17455v1
13. Understanding How Large Language Models Compare Entities, accessed on February 23, 2026, https://www2.informatik.uni-hamburg.de/wtm/publications/2025/LLSW25/2510.16815v1.pdf
14. Learning Shortcuts: On the Misleading Promise of NLU in Language, accessed on February 23, 2026, https://arxiv.org/html/2401.09615v2
15. How Model Architectures Learn and Propagate Bias in Language, accessed on February 23, 2026, https://arxiv.org/html/2505.12381v2
16. How Model Architectures Learn and Propagate Bias in Language, accessed on February 23, 2026, https://www.researchgate.net/publication/391878106_From_n-gram_to_Attention_How_Model_Architectures_Learn_and_Propagate_Bias_in_Language_Modeling
17. Transformers Learn Low Sensitivity Functions: Investigations ... - arXiv, accessed on February 23, 2026, https://arxiv.org/html/2403.06925v2
18. (PDF) Unmasking the Clever Hans effect in AI models - ResearchGate, accessed on February 23, 2026, https://www.researchgate.net/publication/399618464_Unmasking_the_Clever_Hans_effect_in_AI_models_shortcut_learning_spurious_correlations_and_the_path_toward_robust_intelligence
19. Do LLMs Overcome Shortcut Learning? An Evaluation of Shortcut, accessed on February 23, 2026, https://aclanthology.org/2024.emnlp-main.679/
20. Do LLMs Overcome Shortcut Learning? An Evaluation of Shortcut, accessed on February 23, 2026, http://home.ustc.edu.cn/~sa517494/files/emnlp24_yuan.pdf
21. Large Language Models Can be Lazy Learners: Analyze Shortcuts, accessed on February 23, 2026, https://aclanthology.org/2023.findings-acl.284.pdf
22. (PDF) Large Language Models Can be Lazy Learners - ResearchGate, accessed on February 23, 2026, https://www.researchgate.net/publication/372916837_Large_Language_Models_Can_be_Lazy_Learners_Analyze_Shortcuts_in_In-Context_Learning
23. Small Batch Size Training for Language Models: When Vanilla SGD, accessed on February 23, 2026, https://arxiv.org/html/2507.07101v1
24. Small Batch Size Training for Language Models: When Vanilla SGD, accessed on February 23, 2026, https://openreview.net/forum?id=52Ehpe0Lu5&referrer=%5Bthe%20profile%20of%20Andrew%20Gordon%20Wilson%5D(%2Fprofile%3Fid%3D~Andrew_Gordon_Wilson1)
25. publications | Micah Goldblum - GitHub Pages, accessed on February 23, 2026, https://goldblum.github.io/publications/
26. Small Batch Size Training for Language Models: When ... - arXiv, accessed on February 23, 2026, https://arxiv.org/pdf/2507.07101
27. Small Batch Size Training for Language Models: When Vanilla SGD, accessed on February 23, 2026, https://www.researchgate.net/publication/393539348_Small_Batch_Size_Training_

for_Language_Models_When_Vanilla_SGD_Works_and_Why_Gradient_Accumulation_Is_Wasteful
28. Small Batch Size Training for Language Models: When Vanilla SGD, accessed on February 23, 2026, https://arxiv.org/abs/2507.07101
29. COrrect and MItigate Shortcut Learning Behavior in Deep Neural, accessed on February 23, 2026, https://www.researchgate.net/publication/382186152_COMI_COrrect_and_MItigate_Shortcut_Learning_Behavior_in_Deep_Neural_Networks
30. MiMu: mitigating multiple shortcut learning behavior of transformers, accessed on February 23, 2026, https://journal.hep.com.cn/fcs/EN/10.1007/s11704-025-50448-3
31. MiMu: Mitigating Multiple Shortcut Learning Behavior of Transformers, accessed on February 23, 2026, https://arxiv.org/html/2504.10551v1
32. Mitigating Multiple Shortcut Learning Behavior of Transformers, accessed on February 23, 2026, https://www.themoonlight.io/en/review/mimu-mitigating-multiple-shortcut-learning-behavior-of-transformers
33. Mitigating Shortcut Learning via Smart Data Augmentation based on, accessed on February 23, 2026, https://aclanthology.org/2025.coling-main.543/
34. Mitigating Shortcut Learning via Smart Data ... - ACL Anthology, accessed on February 23, 2026, https://aclanthology.org/2025.coling-main.543.pdf
35. Xinyi Sun - ACL Anthology, accessed on February 23, 2026, https://aclanthology.org/people/xinyi-sun/
36. Hongye Tan (谭红叶) - ACL Anthology, accessed on February 23, 2026, https://aclanthology.org/people/hongye-tan/
37. Measuring and Mitigating Post-hoc Rationalization in Reverse, accessed on February 23, 2026, https://arxiv.org/html/2602.14469v1
38. Daily Papers - Hugging Face, accessed on February 23, 2026, https://huggingface.co/papers?q=Rationale-Centric%20Alignment
39. (PDF) Shortcut learning in medical AI hinders generalization, accessed on February 23, 2026, https://www.researchgate.net/publication/380586051_Shortcut_learning_in_medical_AI_hinders_generalization_method_for_estimating_AI_model_generalization_without_external_data
40. Towards Fair Graph Neural Networks via Counterfactual and Balance, accessed on February 23, 2026, https://www.mdpi.com/2078-2489/16/8/704
41. and Inter-modal Causal Attention for Multimodal Sentiment Analysis, accessed on February 23, 2026, https://arxiv.org/html/2508.04999v1